



MANIPAL INSTITUTE OF TECHNOLOGY
MANIPAL
(A constituent unit of MAHE, Manipal)

DSE-3264 - Big Data Analytics Laboratory Manual

| | | | | | |
|---------------------------|---|--|---|---|---|
| Department | : | Data Science Engineering And Computer Applications | | | |
| Course Name & code | : | DSE-3264 & Big Data Analytics Laboratory | | | |
| Semester & branch | : | VI Sem & BTech Data Science & Engineering | | | |
| Name of the faculty | : | Dr. Saraswati Koppad, Dr. Shavantrevva Sangappa Bilakeri | | | |
| No of contact hours/week: | | L | T | P | C |
| | | 0 | 0 | 3 | 1 |

CONTENTS

| Lab No. | Title | Page No |
|---------|---|---------|
| | Course Objectives | |
| | Evaluation Plan | |
| | Instructions to Students | |
| 1 | Understanding Hadoop with HDFS Basic Commands | |
| 2 | Explore advanced HDFS operations | |
| 3 | Map Reduce basics | |
| 4 | Advanced Map reduce | |
| 5 | Exploring Hive | |
| 6 | Exploring Pig | |
| 7 | Exploring HBase | |
| 8 | Spark Basics | |
| 9 | Exploring SparkSQL/SparkML | |
| 10 | Data analytics using Spark. | |
| 11 | Data analytics using Spark | |
| 12 | End Sem exam | |

Course Objectives ·

1. Gain practical experience using big data tools and platforms like Hadoop, Spark, Hive, Pig, and HBase to store, process, and analyze large datasets.
2. Implement Hadoop MapReduce for processing Big Data
3. Apply data analytics techniques to solve problems and extract meaningful insights using big data frameworks.

Course Outcomes:

At the end of this course, students will have the ability to

1. Demonstrate the ability to use big data frameworks such as Hadoop, Spark
2. Apply MapReduce techniques to process large data sets.
3. Demonstrate the ability to use big data tools such as Pig, Hive, and HBase to store and process Big Data
4. Apply analytical methods and techniques to solve data-driven problems.

Evaluation plan : (Tentative)

- Internal Assessment Marks: 60%
- End semester assessment of 2-hour duration: 40 %

Evaluation pattern

| Internal Marks – 60 + End Sem - 40 | |
|------------------------------------|---|
| Internal Marks | Internal Assessment – 40 + Mid-sem - 20 |
| Internal Assessment | Lab observations – 2*11(22) + Viva (18) |
| Lab Observation | Record (1*11) + Execution (1*11) |
| Viva | Quiz and mini-project (18) |

INSTRUCTIONS TO THE STUDENTS

Pre-Lab Session Instructions

- Be on time, adhere to the institution's rules, and maintain decorum.
- Leave your mobile phones, pen drives, and other electronic devices in your bag and keep the bag in the designated place in the lab.
- Must Sign in to the log register provided.
- Make sure to occupy the allotted system and answer the attendance.

In-Lab Session Instructions

- Follow the instructions on the allotted exercises.
- Show the program and results to the instructors on completion of experiments.
- Copy the program and results for the lab record.
- Prescribed textbooks and class notes can be kept ready for reference if required.

General Instructions for the Exercises in Lab

- Academic honesty is required in all your work. You must solve all programming assignments independently, except where group work is authorized. This means you must not take, show, give, or otherwise allow others to take your program code, problem solutions, or other work.
- The programs should meet the following criteria:
 - Programs should be interactive with appropriate prompt messages, error messages if any, and descriptive messages for outputs.
 - Programs should perform input validation (Data type, range error, etc.), give appropriate error messages, and suggest corrective actions.
 - Comments should be used to give the statement of the problem, and every function should indicate the purpose of the function, inputs, and outputs.
 - Statements within the program should be properly indented.
 - Use meaningful names for variables and functions.
 - Make use of constants and type definitions wherever needed.
 - The exercises for each week are divided into three sets:
 - Solved solutions
 - Lab exercises - to be completed during lab hours
 - Additional Exercises - to be completed outside the lab or in the lab to enhance the skill

Questions for lab tests and examinations are not necessarily limited to the questions in the manual but may involve some variations and/or combinations of the questions.

THE STUDENTS SHOULD NOT

- Possess mobile phones or any other electronic gadgets during lab hours.
- Go out of the lab without permission.
- **Change/update any configuration in your allotted system. If so, the student will lose internal assessment marks**

Week 1 - Understanding Hadoop and HDFS Basic Commands

Introduction to the Hadoop Ecosystem and HDFS

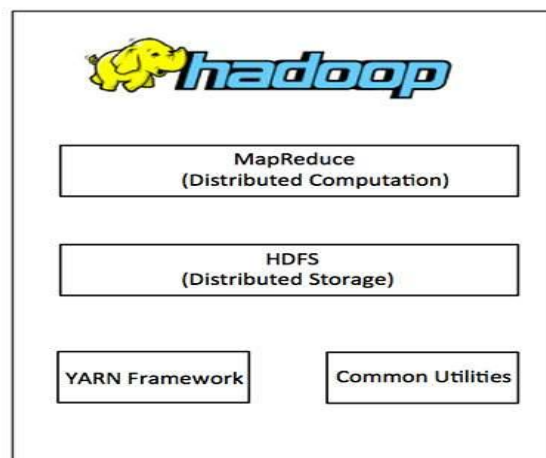
The Hadoop Ecosystem is a collection of tools and frameworks that work together to manage and process big data. HDFS (Hadoop Distributed File System) is the backbone of the Hadoop ecosystem, providing distributed storage and fault tolerance.

Hadoop Ecosystem Overview

The Hadoop Ecosystem is built around the core Hadoop components and includes tools for data storage, processing, querying, and analytics. Here's a breakdown:

Core Components:

1. **HDFS (Hadoop Distributed File System):**
 - Provides distributed storage.
 - Handles large files across multiple machines.
2. **YARN (Yet Another Resource Negotiator):**
 - Manages resources and job scheduling in the cluster.
3. **MapReduce:**
 - Programming model for data processing.



Supporting Tools:

1. **Apache Hive:**
 - Data warehouse tool for querying data in HDFS using SQL-like language (HiveQL).
2. **Apache HBase:**
 - NoSQL database for real-time read/write access to large datasets.
3. **Apache Spark:**
 - Fast, in-memory data processing engine.
4. **Apache Pig:**
 - High-level platform for creating MapReduce programs using a scripting language (Pig Latin).
5. **Apache Sqoop:**
 - Tool for transferring data between Hadoop and relational databases.
6. **Apache Flume:**
 - Tool for collecting and transferring large amounts of log data into HDFS.
7. **ZooKeeper:**
 - Centralized service for managing distributed systems.
8. **Oozie:**
 - Workflow scheduler for Hadoop jobs.

HDFS Basics


HDFS is the storage layer of Hadoop, designed for scalability and fault tolerance. It handles massive data volumes by breaking files into smaller chunks and distributing them across a cluster.

HDFS Features:

- Replication: Ensures fault tolerance by replicating data blocks (default: 3 copies).
- Write-Once, Read-Many: Optimized for batch processing.
- Scalability: Handles petabytes of data across thousands of nodes.

HDFS Architecture:

1. NameNode:
 - Maintains metadata (e.g., file structure, permissions).
 - Coordinates DataNodes but does not store data.
2. DataNodes:
 - Store actual data blocks.
 - Perform read/write operations as instructed by NameNode.
3. Secondary NameNode:
 - Periodically saves snapshots of NameNode metadata (not a backup).



Working with Hadoop:

- Open a new terminal and start the Hadoop service by following commands

```
start-dfs.sh  
start-yarn.sh
```

This command initializes all the required Hadoop daemons for the cluster to become operational.

When you want to stop the services use the command:

```
stop-dfs.sh  
stop-yarn.sh
```

- Check the Status of Hadoop Services (To confirm that all necessary Hadoop services are up and running, you can check their status using the 'jps' command)

```
jps
```

The 'jps' command displays the Java Virtual Machine (JVM) processes and should show a list of Hadoop services running, such as the NameNode, DataNode, ResourceManager, and NodeManager.

```
bdalab@saraswati:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [saraswati]
bdalab@saraswati:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
bdalab@saraswati:~$ jps
3920 NameNode
4996 Jps
4660 NodeManager
4054 DataNode
4520 ResourceManager
4252 SecondaryNameNode
bdalab@saraswati:~$
```

- **Accessing Hadoop Namenode and Resource Manager:**

To access the Hadoop Namenode, open a web browser and enter the following URL:

<http://localhost:9870>

Similarly, to access the Hadoop Resource Manager, open a web browser and enter the following URL:

<http://localhost:8088>

- **Verifying the Hadoop Cluster:**

You can check the Hadoop version by following command

```
hadoop version
```

Create Directories in HDFS:

```
hdfs dfs -mkdir /user/dir_name
```

List Directories in HDFS:

```
hdfs dfs -ls /
hdfs dfs -ls /user/
```

Transfer Files to the Hadoop File System:

```
hdfs dfs -put source-filename /destination_folder/
```


Week-1 Exercise:

1. Practice Linux commands (No need to write in the record book)
 - a. What command would you use to display the current directory?
 - b. How do you list all files, including hidden ones, in a directory?
 - c. Which command is used to create a new directory?
 - d. How can you check the currently logged-in user?
 - e. What does the pwd command do?
 - f. How do you copy a file from one location to another?
 - g. What is the command to move a file?
 - h. How can you delete a directory and all its contents?
 - i. How do you rename a file in Linux?
 - j. What is the difference between `rm` and `rmdir`?
 - k. What command is used to change the permissions of a file?
 - l. How can you view the permissions of a file?
 - m. Which command changes the owner of a file?
 - n. How do you add execute permission for the owner of a file?
 - o. What is the command to display all currently running processes?
 - p. How can you check the IP address of your system?
 - q. How do you display the contents of a file?
 - r. How can you display the current date and time in Linux?
2. Explore Basic HDFS Commands
 - a. Remove directory
 - b. View/Read File Contents
 - c. Download/copy a file from HDFS to the local system
 - d. Copy/move files within HDFS
 - e. Remove file from HDFS
 - f. View file permission